


(a) Image-to-BEV inference

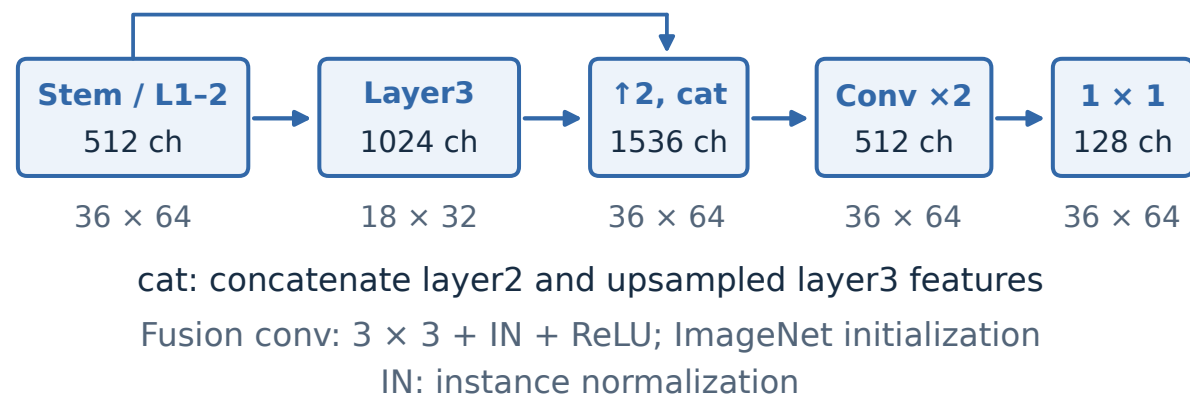
□ Learned module

□ Fixed geometry / fusion



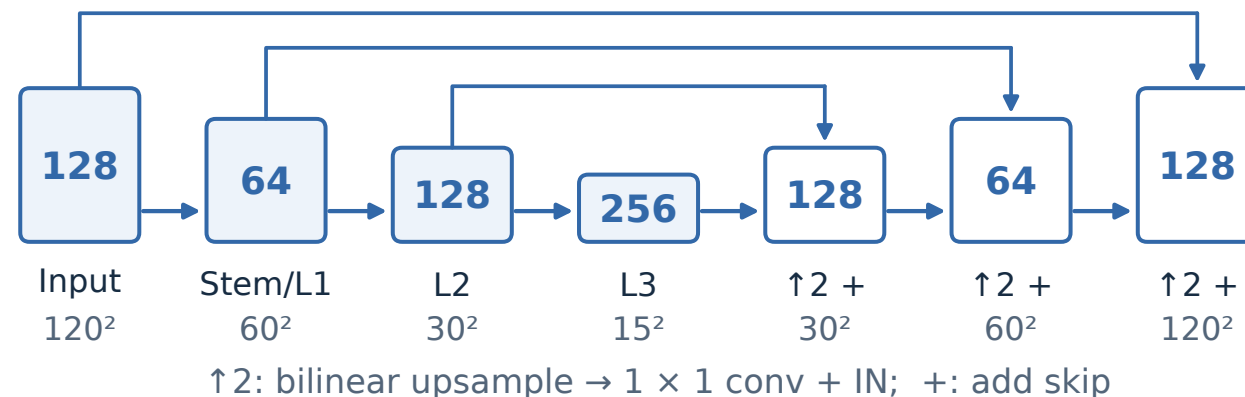
(b) Image encoder

per camera



(c) BEV decoder

$C \times Z \times X$



Batch dimension omitted. Heights are relative to ground; Z/X are BEV axes, Y is height.